



Synchronization

Distributed Systems IT332



Outline

- Clock synchronization
- Logical clocks
- Election algorithms
- Mutual exclusion
- Transactions


Hardware/Software Clocks

- Physical clocks in computers are realized as crystal oscillation counters at the hardware level.
- Usually scaled to approximate physical time t , yielding software clock $C(t)$, $C(t) = \alpha H(t) + \beta$
- $C(t)$ measures time relative to some reference event .
 - Example: 64 bit counter for # of nanoseconds since last boot
- $C(t)$ carries an approximation of real time, never $C(t) = t$.

Hardware/Software Clocks: problems

- Skew: difference between two clocks at one point in time.
- Drift: two clocks tick at different rates.
 - Create ever-widening gap in perceived time.
 - due to physical differences in crystals(quartz oscillators oscillate at slightly different frequencies, plus heat, humidity, voltage, etc.
 - Accumulated drift can lead to significant skew.
 - Clock drift rate: Difference in precision between a prefect reference clock and a physical clock.
 - Usually, 10^{-6} sec/sec, 10^{-7} to 10^{-8} for high precision clocks.


Hardware/Software Clocks



Skew between computer clocks in a distributed system

Clock Synchronization

- Time is unambiguous in centralized systems
 - System clock keeps time, all entities use this for time
- Distributed systems: each node has own system clock
 - Problem: an event that occurred after another event may nevertheless be assigned an earlier time.



Clock Synchronization

➤ Is it possible to synchronize all systems clocks together?

Clock Synchronization

- Clock Synchronization is a mechanism to synchronize the time of all the computers in a DS
- We will study
 - Coordinated Universal Time
 - Clock Synchronization Algorithms
 - Cristian's Algorithm
 - Berkeley Algorithm

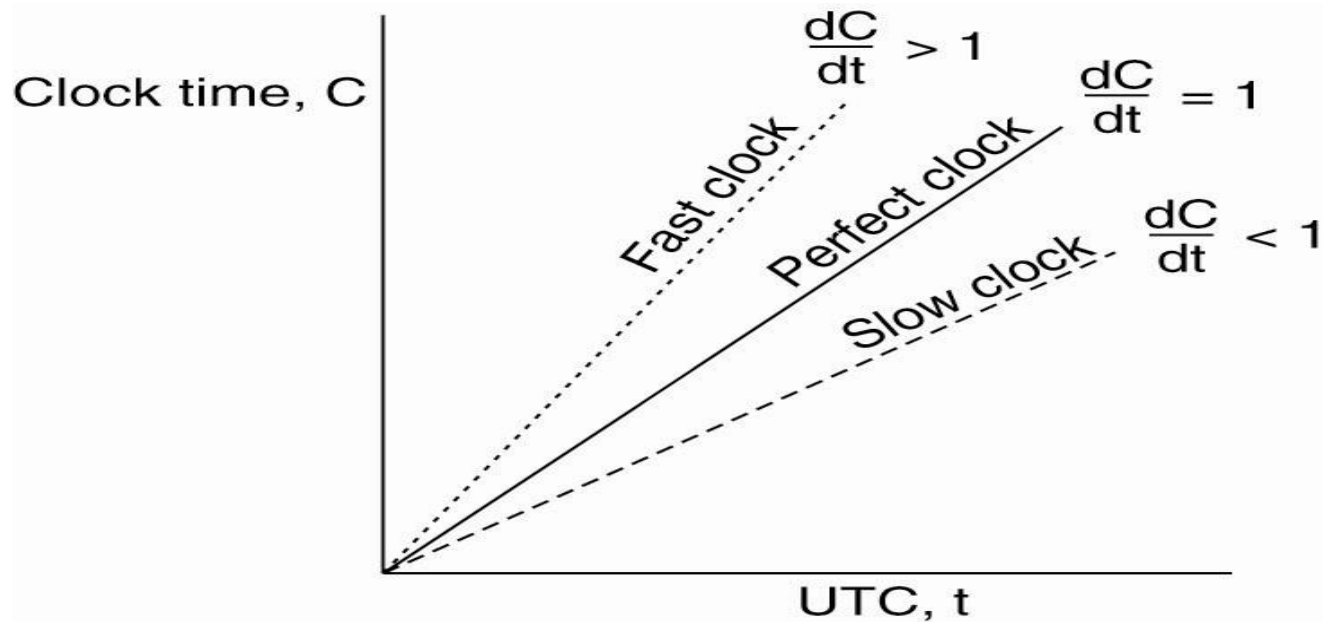
Coordinated Universal Time (UTC)

- All the computers are generally synchronized to a standard time called Coordinated Universal Time (UTC)
 - UTC is the primary time standard by which the world regulates clocks and time. It is available via radio signal, telephone line, satellite (GPS)
- UTC is broadcasted via the satellites
 - UTC broadcasting service provides an accuracy of 0.5 msec
- Computer servers and online services with UTC receivers can be synchronized by satellite broadcasts
 - Many popular synchronization protocols in distributed systems use UTC as a reference time to synchronize clocks of computers

Clock Synchronization

- Need to synchronize machines with a UTC source or with one another
- External synchronization: Synchronize process's clock with an authoritative external reference clock $S(t)$ by limiting skew to a delay bound $\rho > 0$:
 - ➔ $|S(t) - C_i(t)| < \rho$ for all t
- Each clock has a maximum drift rate ρ : $1-\rho \leq dC/dt \leq 1+\rho$
- Internal Synchronization of the local clocks within a distributed system to disagree by not more than a delay bound $\rho > 0$:
 - ➔ $|C_i(t) - C_j(t)| < \rho$ for all i, j, t
- For a system with external synchronization bound of ρ , the internal synchronization is bounded by 2ρ
 - Two clocks may drift by 2ρ in time Δt
 - To limit drift to $\delta \Rightarrow$ resynchronize every $\delta/2\rho$ seconds

Clock Synchronization



The relation between clock time (C) and UTC (t) when clocks tick at different rates.

Getting accurate time

- ❖ Attach GPS receiver to each computer
 ± 1 msec of UTC
- ❖ Attach WWV radio receiver
Obtain time broadcasts from Boulder or DC
 ± 3 msec of UTC (depending on distance)
- ❖ Attach GOES receiver
 ± 0.1 msec of UTC
- ❖ **Not practical solution for every machine**
 - Cost, size, convenience, environment

RPC

- ❖ Synchronize from another machine
 - One with a more accurate clock
- ❖ Machine/service that provides time information:
- ❖ Simplest synchronization technique
 - Issue RPC to obtain time
 - Set time




Does not account for network or processing latency

Cristian's Algorithm

- Use UTC-synchronized time server S
- The time server is passive
- Widely used in LAN.
- Assume that networks delays are symmetric
- Machine A periodically requests time from server B: T₁: request sent and T₄: reply received.
 - A receives time T₂ and T₃ from server, sets clock to T₃+T_{res} where T_{res} is the time to send reply from B to A.
 - Use (T_{req}+T_{res})/2 as an estimate of T_{res}:

$$T_{res} = \frac{(T_2 - T_1) + (T_4 - T_3)}{2}$$




Network Time Protocol (NTP)

- Widely used standard based on Cristian's algorithm
- Improve accuracy by making 8 measurements, take the minimum value for T_{res} as the best estimate for the delay between the two machines.
- A hierarchy of time servers: a time server in level k synchronizes with a server in level $\leq k-1$
 - Level 0 is the atomic clock
- Can we synchronize clocks backward?
 - Clock cannot go backwards: If time needs to be adjusted backward, slow down the clock until time catches up.

Example

- At 5:08:15.100, server B requests time from the time-server A. At 5:08:15.900, server B receives a reply from timeserver A with the timestamp of 5:09:25.300 . (assume there is no processing time at the time-server)
- Send request at 5:08:15.100 (T1)
- Receive response at 5:08:15.900 (T4)
- Response contains 5:09:25.300 (T3)




- What would be the time at server B after synchronization?

Answer

- ❖ Send request at 5:08:15.100 (T_0)
- ❖ Receive response at 5:08:15.900 (T_1)
 - Response contains 5:09:25.300 (T_{server})
- ❖ Elapsed time is $T_1 - T_0$
5:08:15.900 - 5:08:15.100 = 800 msec
- ❖ Best guess: timestamp was generated
400 msec ago
- ❖ Set time to $T_{server} + \text{elapsed time}$
5:09:25.300 + 400 = 5:09:25.700

Berkeley Algorithm

- Keep clocks synchronized with one another
- Assumes that no computer has an accurate time source.
- Machines run Time Daemon.
- One computer is elected as master, others are slaves.
- Master periodically polls slaves for their times and calculate average (including its time)
- Return differences to slaves to synchronize all clocks to average.
- Failure of master => election of a new master



- a) The time daemon asks all the other machines for their clock values
- b) The machines answer
- c) The time daemon tells everyone how to adjust their clock

➤ If two machines do not interact ever. Do we need to synchronize them?

Logical Clocks

- If two machines do not interact, there is no need to synchronize them
- What usually matters is that processes agree on the **order** in which events occur rather than the **time** at which they occurred
 - Absolute time is not important
 - Use **logical clocks**
 - **No concept of happened- when**

Event Ordering

- **Problem:** define a global ordering of all events that occur in a system
- Events in a single process or machine are locally ordered
- In a distributed system:
 - No global clock, local clocks may be unsynchronized
 - Can not order events on different machines using local times
- **Key idea [Lamport]**
 - Processes exchange messages
 - Message must be sent before received
 - Send/receive used to order events and synchronize logical clocks


Happened-Before (HB) Relation


- If A and B are events in the same process and A occurs before B, then $A \rightarrow B$
- If A represents sending of a message and B is the receipt of this message, then $A \rightarrow B$ ($\text{clock}(A) < \text{clock}(B)$)
- Relation is transitive:
 - $A \rightarrow B$ and $B \rightarrow C$ implies $A \rightarrow C$
- Unordered events are concurrent
 - $A \nrightarrow B$ and $B \nrightarrow A$ implies $A \parallel B$

Lamport's Logical Clocks

- **Goal:** assign timestamps to events such that
If $A \rightarrow B$ then $\text{timestamp}(A) < \text{timestamp}(B)$
- Lamport's Algorithm
 - Each process i maintains a logical clock L_i
 - Whenever an event occurs locally at i , $L_i = L_i + 1$
 - When i sends message to j , piggyback L_i
 - When j receives message from i , $L_j = \max(L_i, L_j) + 1$
- In this algorithm, $A \rightarrow B$ implies $L(A) < L(B)$, but $L(A) < L(B)$ does not necessarily imply $A \rightarrow B$

Lamport's Logical Clocks: An Example



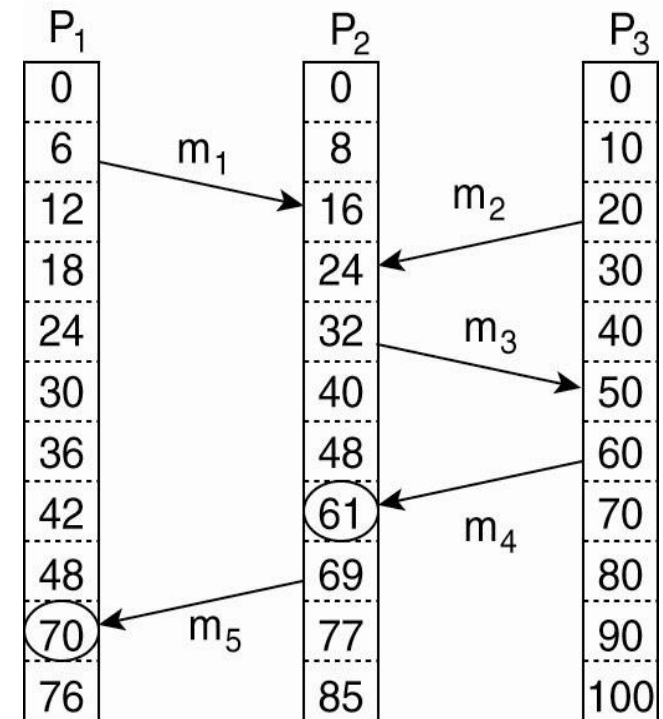


(a) Three processes, each with its own clock. The clocks run at different rates.

(b) Lamport's algorithm corrects the clocks.

Causality

- Lamport's logical clocks:
 - If $A \rightarrow B$ then $L(A) < L(B)$
 - Problem: Reverse is not true!!
 - Nothing can be said about events by comparing timestamps!
- Need to capture causality
 - If $A \rightarrow B$ then A causally precedes B
 - Need a timestamping mechanism such that:
 - $T(A) < T(B)$ iff A causally precedes B




Event A: m_1 is received at $t=16$
 Event B: m_2 is sent at $t=20$
 $L(A) < L(B)$, but A does not causally precede B.

Solution: Vector Clocks

- Each process i maintains a vector clock V_i of size N , where N is the number of processes
 - $V_i[i]$ = number of events that have occurred at process i
 - $V_i[j]$ = number of events i knows have occurred at process j
 - Initially, $V_i[j] = 0$ for all $i, j = 1, \dots, N$
- Update vector clocks as follows:
 - Local event at p_i : **increment $V_i[i]$ by 1 (before time stamping event)**
 - When a message is sent from p_i to p_j : **piggyback vector V_i**
 - Receipt of a message from p_i by p_j : **$V_j[k] = \max(V_j[k], V_i[k])$, $j \neq k$; $V_j[j] = V_j[j] + 1$**
 Receiver is told about how many events the sender knows occurred at another process k
- We have $V(A) < V(B)$ iff A causally precedes B !
 - $V(A) < V(B)$ iff for all i , $V(A)[i] \leq V(B)[i]$ and there exists k such that $V(A)[k] < V(B)[k]$
- A and B are concurrent iff $V(A) \nless V(B)$ and $V(B) \nless V(A)$

Vector Clock: An Example



Election Algorithms

- Many distributed algorithms need one process to act as a leader or coordinator
 - How to select this process dynamically
 - Doesn't matter which process does the job, just need to pick one
 - Example: pick a master in Berkeley clock synchronization algorithm
- Election algorithms: technique to pick a unique coordinator
 - Assumption: each process has a unique ID
 - Goal: find the non-crashed process with the highest ID

Bully Algorithm

➤ Assumptions

- Each process knows the ID and address of every other process
- Communication is reliable

➤ A process initiates an election if it just recovered from failure or it notices that the coordinator has failed


➤ Three types of messages: Election, OK, Coordinator

- Several processes can initiate an election simultaneously
 - Need consistent result

Bully Algorithm Details

- Any process P can initiate an election
- P sends Election messages to all process with higher IDs and awaits OK messages
 - If no OK messages, P becomes coordinator and sends Coordinator messages to all processes with lower IDs
 - If it receives an OK, it drops out and waits for an Coordinator message
- If a process receives an Election message
 - Immediately sends Coordinator message if it is the process with highest ID
 - Otherwise, returns an OK and starts an election
- If a process receives a Coordinator message, it treats sender as the coordinator

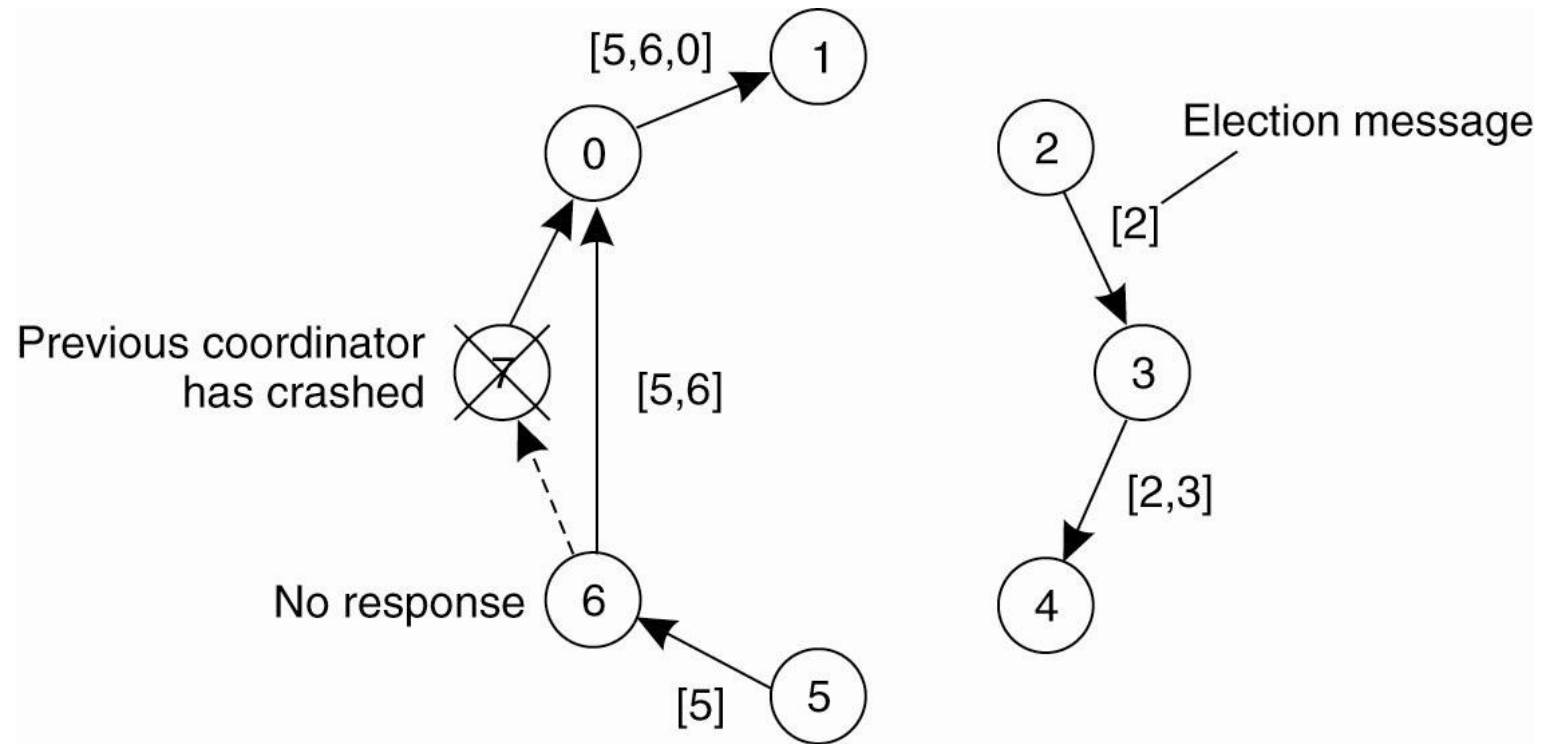
Bully Algorithm Example



Ring Algorithm

- Processes are arranged in a logical ring, each process knows the structure of the ring
- A process initiates an election if it just recovered from failure or it notices that the coordinator has failed
- Initiator sends Election message to closest downstream node that is alive
 - Election message is forwarded around the ring
 - Each process adds its own ID to the Election message
- When Election message comes back, initiator picks node with highest ID and sends a Coordinator message specifying the winner of the election
 - Coordinator message is removed when it has circulated once.
- Multiple elections can be in progress

Ring Algorithm Example





Comparison of Bully and Ring Algorithms

- Assume n processes and one election in progress
- Bully algorithm
 - Worst case: initiator is node with lowest ID
 - Triggers $n-2$ elections at higher ranked nodes: $O(n^2)$ messages
 - Best case: initiator is node with highest ID
 - Immediate election: $n-1$ messages
- Ring algorithm
 - $2n$ messages always


Election in Wireless Environments

- Goal: elect the best leader (e.g., node with longest battery lifetime)



Node a initiates an election.

Election in Wireless Networks



In the end, source a notes that h is the best leader and broadcasts this info to all nodes.

Mutual Exclusion

- Processes in a distributed system may need to simultaneously access the same resource
- Need to grant mutual exclusive access to shared resources by processes
- Solutions:
 - Via a centralized server (Centralized algorithm)
 - Decentralized, using a peer-to-peer system (Decentralized algorithm)
 - Distributed, with no topology imposed (Distributed algorithm)
 - Distributed, along a logical ring (A token ring algorithm)



Terminology

- ❖ In concurrent programming a **critical section** is a piece of code that accesses a shared resource that must not be concurrently accessed by more than one thread of execution.
- ❖ **Mutual exclusion** (ME, often abbreviated to mutex) algorithms are used in concurrent programming to avoid the simultaneous use of a common resource, such as a global variable, by pieces of computer code called critical sections.




Mutual Exclusion

- ❖ Prevent simultaneous access to a resource
- ❖ Two basic kinds:
 - **Permission based**
 - A Centralized Algorithm
 - A Decentralized Algorithm
 - A Distributed Algorithm
 - **Token based**
 - A Token Ring Algorithm


Centralized Mutual Exclusion

- Assume processes are numbered
- One process is elected coordinator
- Every process needs to check with coordinator before entering the critical section
- To obtain exclusive access: send request, await reply
- To release: send release message
- Coordinator:
 - Receive request: if resource is available and queue empty, sendOK; if not, queue request
 - Receive release: remove next request from queue and sendOK


Centralized Mutual Exclusion



(a)



(b)



(c)

- a) Process 1 asks the coordinator for permission to access a shared resource. Permission is granted.
- b) Process 2 then asks permission to access the same resource. The coordinator does not reply.
- c) When process 1 releases the resource, it tells the coordinator, which then replies to 2.


Properties of Centralized Mutual Exclusion

- Simulate centralized locking with blocking calls
- Advantages
 - Fair: requests are granted in the order they were received
 - Simple: three messages per use of a resource (request, OK, release)
 - No starvation
- Drawbacks:
 - Single point of failure
 - Performance bottleneck in large distributed systems
 - How do you detect a dead coordinator?
 - A process can not distinguish between “permission denied” from a dead coordinator – No response from coordinator in either case

A Decentralized Algorithm

- Each resource is replicated n times. Each replica has its own coordinator
- Access requires majority vote from $m > n/2$ coordinators.
 - Nonblocking: coordinators return OK or “no”
- Coordinator crashes => forgets previous votes (i.e., resets itself)
- If request is denied, process will back off for a randomly-chosen time, and try again
- **Drawbacks**
 - Low resource utilization when many nodes want to access the same resource
 - Starvation can occur


Example




A Distributed Algorithm

- Based on event ordering and time stamps
 - Each process maintains a logical clock
- Process k enters critical section as follows
 - Increment logical clock: $L_k = L_k + 1$
 - Multicast a message (L_k, k) to all other processes
 - Wait until a reply is received from every other process
 - Enter critical section
- Upon receiving a request message, process j
 - Sends an OK message if outside of critical section
 - If already in critical section, does not reply, queue the request
 - If wants to enter critical section, sends an OK message if $(L_k, k) < (L_j, j)$, else queue the request
- When a process is finished with the critical section
 - Send OK messages to all processes on its queue and delete them from the queue


A Distributed Algorithm



(a)



(b)



(c)

a) Two processes want to access a shared resource at the same moment.

b) Process 0 has the lowest timestamp, so it wins.



c) When process 0 is done, it sends OK to 2, so 2 can now go ahead

Properties of Distributed Algorithm

- Fully distributed
- N points of failure!
- All processes are involved in all decisions
 - Any overloaded process can become a bottleneck
- Improvements
 - Shows that a fully distributed system is possible.
 - When a request comes in, always sends a reply granting or denying permission.
 - This helps detect dead processes
 - Enter critical section when the process has got permission from a simple majority of the other processes

A Token Ring Algorithm

- Assume known group of processes
 - Some order can be imposed
 - Construct logical ring in software
 - Process communicates with neighbour
- Use a token to arbitrate access to critical section
- P_0 gets token to Resource R
- Token circulates around the ring: P_i passes to $P_{i+1} \bmod N$
- Must wait for token before entering CS
- Process which acquires a token checks if it needs the critical section
 - If No: Pass the token to neighbor
 - If yes: access resource, holds token until done



Features

- Only one process at a time has a token
 - ME is guaranteed
- Order well defined
 - No starvation!
- If token is lost (e.g:Process died)
 - It will be have regenerated
- Detecting token loss is non-trivial
- Doesn't guarantee FIFO order (sometimes it is undesirable)
- Failing nodes can break the ring

Comparison

| Algorithm | Delay before entry | Messages per entry / exit | Problems |
|---------------|--------------------|---------------------------|--|
| Centralized | 2 | 3 | Coordinator crashes |
| Decentralized | $2mk$ | $3mk$ | Starvation, Low efficiency |
| Distributed | $2(n-1)$ | $2(n-1)$ | Crash of any process |
| Token Ring | 0 to $n-1$ | 1 to infinity | Token may be lost, Ring can be broken if processes crash |